

Introduction to System Programming


System Calls For Semaphore Management

Outline

- ❑ System Calls for Semaphore Management
 - ❑ semget
 - ❑ semctl
 - ❑ semop
 - ❑ fsync

Semaphore structure

Sub-Semaphores



semget System Call

- ❑ Semget Function : To create a semaphore, or gain access to one that exists.
- ❑ Include: `<sys/types.h>` `<sys/ipc.h>` `<sys/sem.h>`
- ❑ Command: ***int semget(key_t key, int nsems, int semflg);***

Arguments

- ***key_t key***: used to identify a semaphore set
- ***int nsems***: the number of semaphores in the set.
- ***int semflg***: specify access permissions and/or special creation condition(s).

- ❑ Returns: Success: the semaphore identifier (semid);

Failure :-1; Sets errno: yes

- ❑ To see whether semaphore is created use :

\$ **ipcs -s**

IDKey mode Owner nsems

Creating and Accessing Semaphore Sets

```
main()
{
    key=(key_t)0x20;
    nsem=1
    semid=semget(key, nsem, IPC_CREAT|0666)
}
```



Read-alter mode

Flag: IPC_EXCL: Exclusive creation of semaphore

IPC_CREAT|0666|IPC_EXCL

Program to create semaphore (S1.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
void main()
{
    int semid,key,nsem;
    key=(key_t)0x20;
    nsem=0;
    semid=semget(key,nsem,IPC_CREAT|0666);
    printf("created semaphore with id :%d\n",semid);
}
```

\$. /a.out
created semaphore
with id :-1

To see whether semaphore is created or not

//command to verify the creation of semaphore

\$ipcs -s

----- Semaphore Arrays -----

key	semid	owner	perms	nsems
-----	-------	-------	-------	-------

Program to create semaphore (S2.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
void main()
{
    int semid,key,flag,nsem;
    key=(key_t)0x20;
    flag=IPC_CREAT|0666;
    nsem=1;
    semid=semget(key,nsem,flag);
    printf("created semaphore with id :%d\n",semid);
}
```


Program to create semaphore (S2.txt)

```
$/a.out
```

```
created semaphore with id :0
```

```
$ipcs -s
```

```
----- Semaphore Arrays -----
```

key	semid	owner	perms	nsems
0x000000020	0	student	666	1



To find maximum number of semaphore sets available in linux system (s3.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<stdlib.h>
void main()
{
    int semid,nsemset,key,flag,nsem;
    nsem=1;
    flag=IPC_CREAT|0666;
    for(nsemset=0;;nsemset++)
    {
        key=(key_t)nsemset;
        semid=semget(nsemset,nsem,flag);
```



To find maximum number of semaphore sets available in linux system (s3.c)

```
if(semid > 0)
    printf("created semaphore with id :%d\n",semid);
else
{
    printf("Maximum number of semaphore set are  %d
           \n",nsemset);
    exit(0);
}
}
```

To find maximum number of semaphore sets available in linux system (S3.txt)

```
$/a.out
```

```
created semaphore with id :32769
```

```
created semaphore with id :65538
```

```
created semaphore with id :98307
```

```
created semaphore with id :131076
```

```
created semaphore with id :163845
```

```
created semaphore with id :196614
```

```
created semaphore with id :229383
```

```
created semaphore with id :262152
```

```
created semaphore with id :294921
```

```
created semaphore with id :327690
```

```
created semaphore with id :360459
```

```
created semaphore with id :393228
```

```
created semaphore with id :425997
```

To find maximum number of semaphore sets available in linux system (S3.txt)

created semaphore with id :458766
created semaphore with id :491535
created semaphore with id :524304
created semaphore with id :557073
created semaphore with id :589842
created semaphore with id :622611
created semaphore with id :655380
created semaphore with id :688149
created semaphore with id :720918
created semaphore with id :753687
created semaphore with id :786456
created semaphore with id :819225
created semaphore with id :851994
created semaphore with id :884763

To find maximum number of semaphore sets available in linux system (S3.txt)

created semaphore with id :917532
created semaphore with id :950301
created semaphore with id :983070
created semaphore with id :1015839
created semaphore with id :1048608
Maximum number of semaphore set are 32

When issued : (\$ipcs -s)

----- Semaphore Arrays -----

<u>key</u>	<u>semid</u>	owner	perms	<u>nsems</u>
0x000000020	0	student	666	1
0x000000000	32769	student	666	1
0x000000001	65538	student	666	1
0x000000004	163845	student	666	1
0x000000005	196614	student	666	1
0x000000002	98307	student	666	1
0x000000003	131076	student	666	1
0x000000006	229383	student	666	1
0x000000007	262152	student	666	1
0x000000008	294921	student	666	1
0x000000009	327690	student	666	1
0x00000000a	360459	student	666	1
0x00000000b	393228	student	666	1
0x00000000c	425997	student	666	1

When issued : (\$ipcs -s)

0x00000000d	458766	student	666	1
0x00000000e	491535	student	666	1
0x00000000f	524304	student	666	1
0x000000010	557073	student	666	1
0x000000011	589842	student	666	1
0x000000012	622611	student	666	1
0x000000013	655380	student	666	1
0x000000014	688149	student	666	1
0x000000015	720918	student	666	1
0x000000016	753687	student	666	1
0x000000017	786456	student	666	1
0x000000018	819225	student	666	1
0x000000019	851994	student	666	1
0x00000001a	884763	student	666	1
0x00000001b	917532	student	666	1
0x00000001c	950301	student	666	1
0x00000001d	983070	student	666	1
0x00000001e	1015839	student	666	1
0x00000001f	1048608	student	666	1

To find maximum number of semaphore in each set

(s4.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<stdlib.h>
void main()
{
    int semid,i,key,flag,nsem;
    key=(key_t)0x30;
    nsem=1;
    flag=IPC_CREAT|0666;
    for(i=0;;i++)
    {
        nsem=i+1;
        semid=semget(key,nsem,flag);
```

To find maximum number of semaphore in each set

(s4.c)

```
if(semid > 0)
    printf("created semaphore with id :%d\n",semid);
else
{
    printf("Maximum number of semaphore set are %d\n",i);
    exit(0);
}
semctl(semid,0,IPC_RMID,0);
}
}
```

To find maximum number of semaphore in each set (S4.txt)

\$/a.out

created semaphore with id :1081377

created semaphore with id :1114145

created semaphore with id :1146913

created semaphore with id :1179681

created semaphore with id :1212449

created semaphore with id :1245217

● ● ●

Maximum number of semaphore set are 32000

Semaphore exclusivity (s7.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<stdlib.h>
void main()
{
    int semid,key,flag,nsem,i;
    for(i=0;i<2;i++)
    {
        key = (key_t)0x30;
        flag=IPC_CREAT|0666|IPC_EXCL;
        nsem=1;
```

Semaphore exclusivity (s7.c)

```
semid=semget(key,nsem,flag);
if(semid > 0)
    printf("created semaphore with id :%d \n",
           semid);
else
    perror("semget");
}
```

Semaphore exclusivity (s7.txt)

❑ **Note**: delete the semaphore with KEY 30 before running the above program.

❑ To do this

- **\$ipcrm -S 0x30**
- **\$/a.out**

created semaphore with id :999424031

***semctl* System Call - Semaphore Control**

- ❑ Function :To perform a variety of generalized control operations on the system semaphore structure, on the semaphores as a set and on individual semaphores.
- ❑ Include **<sys/types.h> <sys/ipc.h> <sys/sem.h>**
- ❑ Command: ***int semctl (int semid, int semum, int cmd, /* union semun arg*/ ...);***
- ❑ **Returns :** **Success :** 0 or the value requested
Failure : -1
Seterrno : - Yes

Arguments

- ***int semid***: a valid semaphore identifier.
- ***int semum***: the number of semaphores in the semaphore set.
- ***int cmd***: an integer command value (IPC_STAT, IPC_SET, ..).
- ***arg***: union of type ***semun***.

***semctl* System Call - Semaphore Control**

- ❑ The **semctl** system call takes **four** arguments:
- The first argument, `semid`, is a valid semaphore identifier that was returned by a previous `semget` system call.
 - The second argument, `semnum`, is the number of semaphores in the semaphore set.
 - The third argument to `semctl`, `cmd`, is an integer command value. The `cmd` value directs `semctl` to take one of several control actions. Each action requires specific access permissions to the semaphore control structure.

cmd values of *semctl* Call

- ❑ IPC_STAT: return the current values of the `semid_ds` structure for the indicated semaphore identifier.
 - ❑ IPC_SET: modify a restricted number of members in the `semid_ds` structure.
 - ❑ IPC_RMID: remove the semaphore set.
 - ❑ GETALL: return the current values of the semaphore set.
 - ❑ SETALL: Initialize all semaphores in a set to the values stored in the array referenced by the fourth arguments to `semctl`.
 - ❑ GETVAL: return the current of the individual semaphore referenced by the value of the ***semnum*** argument.
 - ❑ SETVAL: set the value of a single semaphore in a set
-

***semctl* System Call - Semaphore Control**

- ❑ The **semctl** system call takes **four** arguments:
- ❑ The fourth argument to **semctl**, **arg**, is a union of type **semun**. Given the action specified by the preceding **cmd** argument, the data in **arg** can be one of any of the following four values:
 - An integer used with **SETVAL** to indicate a specific value for a particular semaphore within the semaphore set.
 - A reference to a **semid_ds** structure where information is returned when **IPC_STAT** or **IPC_SET** is specified.
 - A reference to an array of type unsigned short integers; the array is used either to initialize the semaphore set or as a return location when specifying **GETALL**.
 - A reference to a **seminfo** structure when **IPC_INFO** is requested.

union *semun* in the *semctl* call

A ***union*** is a later-day version of the Pascal *variant record*. It is a data structure that can take on multiple forms.

semctl() requires a union to handle the different kinds of data that can be provided to it or received from it.

```
union semun {  
    int      val;  
    struct    semid_ds *buf;  
    ushort    * array;  
} arg; // declares a semun named arg
```

The value in *arg* is one of:

- int *val*: an integer (0 or others),
- struct *semid_ds* **buf*: a reference to a *semid_ds* structure,
- ushort **array*: the base address of an array of short integers (the values for the semaphore(s)).

Getting and setting Semaphore values (s8.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
#include<stdlib.h>
#include<errno.h>
void main()
{
    int semid,retval;
    semid=semget(0x20,1,0666|IPC_CREAT);
    retval=semctl(semid,0,GETVAL,0);
    printf("value returned is %d\n",retval);
}
```

- \$./a.out

value returned is 0

Setting different value to Semaphore (s9.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
#include<stdlib.h>
#include<errno.h>
```

Setting different value to Semaphore (s9.c)

```
void main()
{
    int semid,retval;
    semid=semget(0x20,1,0666|IPC_CREAT);
    semctl(semid,0,SETVAL,1);
    retval=semctl(semid,0,GETVAL,0);
    printf("value of the semaphore after setting is %d\n",retval);
    semctl(semid,0,SETVAL,2);
    retval=semctl(semid,0,GETVAL,0);
    printf("value of the semaphore after setting is %d\n",retval);
}
```

- **\$. /a.out**

value of the semaphore after setting is 1

value of the semaphore after setting is 2

Who is using Resources (s10.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
#include<errno.h>
void main()
{
    int semid,retval;
    semid=semget(0x20,1,0666|IPC_CREAT);
    semctl(semid,0,SETVAL,1);
    retval=semctl(semid,0,GETPID,0);
    printf("PID returned by semctl is %d and actual PID is
%d\n",retval,getpid());
```

- **\$/a.out**
- PID returned by semctl is 3409 and actual PID is 3409
- Note kill the semaphore first and then run this program twice

Who is using Resources . Here subsemaphore is intialized (s11.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
#include<errno.h>
void main()
{
```

- Note : remove the semaphore before running the program
- **\$/a.out**
- PID returned by semctl is 3409 and actual PID is 3429

```
    int semid,retval;
    semid=semget(0x20,1,0666|IPC_CREAT);
    retval=semctl(semid,0,GETPID,0);
    printf("PID returned by semctl is %d and actual PID is
           %d\n",retval,getpid());
    retval=semctl(semid,0,SETVAL,1);
```

```
}
```


Sem_id

S

Semaphore structure

/* One sem_array data structure for each set of semaphores in the system. */

```
struct sem_array {  
    struct kern_ipc_perm    sem_perm; /* permissions .. see ipc.h */  
    time_t                  sem_otime; /* last semop time */  
    time_t                  sem_ctime; /* last change time */  
    struct sem              *sem_base; /* ptr to first semaphore in array */  
    struct sem_queue        *sem_pending; /* pending operations to be processed */  
    struct sem_queue        **sem_pending_last; /* last pending operation */  
    struct sem_undo          *undo; /* undo requests on this array */  
    unsigned long           sem_nsems; /* no. of semaphores in array */  
};
```

semid_ds

```
struct ipc_perm
{
    key_t  key;
    ushort uid; /* owner euid and egid */
    ushort gid;
    ushort cuid; /* creator euid and egid */
    ushort cgid;
    ushort mode; /* access modes see mode flags below */
    ushort seq; /* slot usage sequence number */
};
```

```
struct sem {
```

```
    u_short semval;
```

```
    short  sempid;
```

```
    u_short semncnt;
```

```
    u_short semzcnt;
```

```
};
```

Waiting for positive

Waiting for zero

Sembuf

```
struct sembuf
```

```
{  
    ushort sem_num;    → Sub semaphore
```

```
    short  sem_op;
```

```
    short sem_flg;
```

```
}    → 0, IPC_NOWAIT, SEM_UNDO
```

```
int semop(int semid, struct sembuf *sops,  
unsigned nsops);
```

***semop* Call - Semaphore Operation**

- ❑ Function : to perform operations on individual semaphores.
- ❑ Include: **<sys/types.h> <sys/ipc.h> <sys/sem.h>**
- ❑ Command: ***int semop (int semid, struct sembuf *sops, size_t nsops);***
- ❑ Returns: Success: 0; Failure; -1; Sets errno: Yes.
- ❑ Arguments
 - ❑ ***int semid***: semaphore identifier.
 - ❑ ***struct sembuf *sops***: a reference to the address of an array of semaphore operations that will be performed on the semaphore set denoted by the semid value.
 - ❑ ***size_t nsops***: the number of elements in the array of semaphore operations.

Actions Taken by *semop*

□ If ***semop*** value:

- is positive: : Add ***sem_op*** to ***semval***. This is a release of a resource
- is zero: : The caller will block until the semaphore's value becomes zero.
- is negative: The caller is blocked until the semaphore's value (***semval***) becomes greater than or equal to the absolute value of ***sem_op***. Then, the absolute value of ***sem_op*** *is subtracted* from ***semval***.

SEMOP semaphore (y.c)

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
void main()
{
    int semid,pid;
    struct sembuf sop;
    semid=semget(0x20,1,0666|IPC_CREAT);
    pid = fork();
```

SEMOP semaphore (y.c)

```
if(pid == 0)
{
    sleep(2);
    printf("Child before semop\n");
    sop.sem_num = 0;
    sop.sem_op = 0;
    sop.sem_flg = 0;
    semop(semid,&sop,1);
    printf("child over\n");
}
```


SEMOP semaphore (y.c)

else

{

printf("parent before 1st semctl\n");

semctl(semid,0,SETVAL,1);

printf("Parent sleeping\n");

sleep(5);

printf("parent before 2nd semctl\n");

semctl(semid,0,SETVAL,0);

printf("Parent over\n");

}

}

Running the previous program

- parent before 1st semctl
Parent sleeping
parent before 2nd semctl
Parent over
Child before semop
child over

----- Semaphore Arrays -----

key	semid	owner	perms	nsems
0x00000020	32768	student	666	1

Sturcture of semaphore

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
void main()
{
    int semid;
    struct semid_ds status;
    semid=semget((key_t)0x20,10,0666|IPC_CREAT);
    semctl(semid,0,IPC_STAT,&status);
    printf("No of semphoares in set are %ld\n",status.sem_nsems);
    printf("My user id is %u\n",getuid());
    printf("Owner user id is %u\n",status.sem_perm.uid);
    printf("My group id is %u\n",getgid());
```

Sturcture of semaphore

```
printf("Owner group id is %u\n",status.sem_perm.gid);  
printf("Creator user id is %u\n",status.sem_perm.cuid);  
printf("Creator group id is %u\n",status.sem_perm.cgid);  
printf("Access mode is %o\n",status.sem_perm.mode);  
}
```

No of semphoares in set are 94610192697616

My user id is 1001

Owner user id is 0

My group id is 1001

Owner group id is 1970169159

Creator user id is 0

Creator group id is 9

Access mode is 0

fsync, fdatasync - synchronize a file's in-core state with storage device

- ❑ **fsync()** transfers ("flushes") all modified in-core data of (i.e., modified buffer cache pages for) the file referred to by the file descriptor *fd* to the disk device (or other permanent storage device) where that file resides.
- ❑ The call blocks until the device reports that the transfer has completed.
- ❑ It also flushes metadata information associated with the file (see **stat(2)**).
- ❑ Calling **fsync()** does not necessarily ensure that the entry in the directory containing the file has also reached disk.

fsync, fdatasync - synchronize a file's in-core state with storage device

- ❑ Include : **<unistd.h>**
- ❑ Command : **int fsync(int *fd*);**
int fdatasync(int *fd*);
- ❑ Return : On success, zero is returned.
On error, -1 is returned,
and *errno* is set appropriately.

Program for fsync()

```
#include <stdio.h>
#include <fcntl.h>
int main()
{
    char my_write_str[] = "1234567890";
    char my_read_str[100];
    char my_filename[] = "m1.txt";
    int my_file_descriptor, close_err;
    /* Open the file. */
    my_file_descriptor = open (my_filename, O_RDWR |
    O_CREAT | O_TRUNC);
```

Program for fsync()

```
/* Write 10 bytes of data and make sure it's written */  
write (my_file_descriptor, (void *) my_write_str, 10);  
fsync (my_file_descriptor);
```

```
/* Seek the beginning of the file */  
lseek (my_file_descriptor, 0, SEEK_SET);
```

```
/* Read 10 bytes of data */  
read (my_file_descriptor, (void *) my_read_str, 10);  
/* Terminate the data we've read with a null character  
*/  
my_read_str[10] = '\0';
```


Program for fsync()

```
printf ("String read = %s.\n", my_read_str);  
close (my_file_descriptor);  
return 0;  
}
```

